


Evaluation of LLMs and RAG Pipelines in a Smart Way



by Anurag Pathak



Why Evaluation Matters

1 Manual Checks Insufficient


Unscalable and prone to human error

2 Need Scalable, Reliable Evaluation

Ensure consistent performance

3 Impact on Production

Affects model alignment and user trust



Qualitative vs Quantitative Metrics

Qualitative Metrics

Human-centric assessment

- Interpretability
- Coherence
- Bias detection

Quantitative Metrics

Numeric, formula-based

- Accuracy
- F1 Score
- Perplexity

Confusion Matrix

True Positive (TP)

Correctly identified positive

False Positive (FP)

Incorrectly identified positive

False Negative (FN)

Incorrectly identified negative

True Negative (TN)

Correctly identified negative

n=165	Predicted: NO	Predicted: YES	
Actual: NO	TN = 50	FP = 10	60
Actual: YES	FN = 5	TP = 100	105
	55	110	



Automated Evaluation Metrics



Text Overlap

BLEU, ROUGE,
METEOR



Classification

Accuracy, Precision,
Recall, F1



Generative Models


Perplexity, Exact
Match



Ethics/Alignment

Toxicity, Bias detection

Human Evaluation Metrics



Testing LLM Inference



Prompt Engineering

Crafting effective inputs



Zero/Few-shot

Evaluating adaptation



Benchmarks

HELM, MMLU, TruthfulQA



Performance

Latency and throughput

RAG Pipeline Evaluation



Retriever Metrics

Recall@k, Context Relevancy, Context Precision



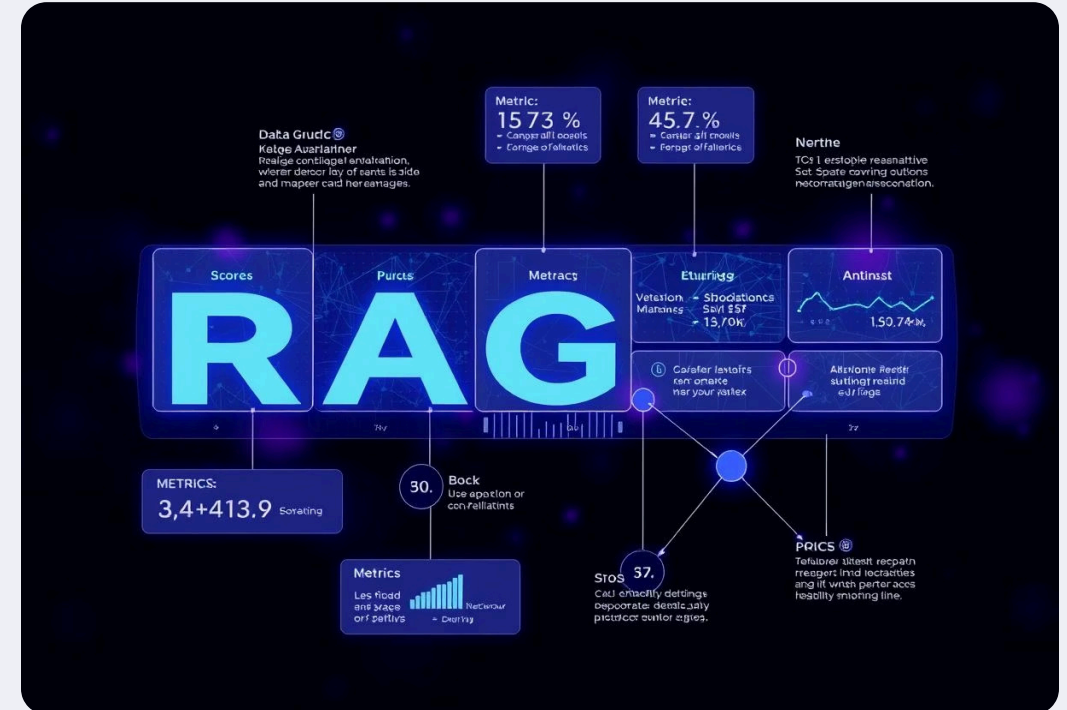
Generator Metrics

Factual grounding, hallucination rate, bias, answer correctness



End-to-End Testing

Crucial for overall quality



Evaluation Frameworks & Tooling

1

Giskard Introduction

Leading evaluation framework



Bias & Robustness

Automated checks



Human Feedback

Seamless integration



CI/CD Integration

For LLM pipelines

Best Practices & Pitfalls

Hybrid Evaluation

Combine automated and human insights

Beyond BLEU/ROUGE

Use diverse metrics for depth

Update Datasets

Keep evaluation data fresh

Align with Goals

Match eval to user needs

